FDE Labs Challenge

"The Data Room" · 19 October – 13 November 2026

Four weeks, as a side project around your normal work, to take a real business process that is cleared by hand today and reinvent it as a working agentic AI system — and prove the improvement with KPIs. Your laptop, your AI tooling, your architecture.

The task

Atlas Industrial acquired a competitor and inherited a chaotic data room of vendor contracts — PDFs, Word docs, a low-quality scan, a forwarded email and a legacy CSV. Analysts clear it by hand: open each document, hand-key fields into the ERP, eyeball risk, email Legal for exceptions. Slow, error-prone, unauditable.

Replace it. For every document your system must:

- Take **exactly one action** — load, escalate or quarantine — via the mock ERP API, with a logged rationale behind every decision.
- Produce records that **conform to the schema**. The API returns 422 on anything that doesn't.
- Apply the client's **policy** correctly, including the data-governance fields and the dual-approval guardrail.
- **Escalate rather than guess** when the data is ambiguous, low-confidence, conflicting or manipulative.

The inputs are messy, the output format is strict, the policy is non-negotiable and the integration is live. **There is no starter code** — the architecture is entirely yours.

What you submit

- **Your code**, zipped, with a README that runs it from clean
- **A 5–10 min recorded demo** — end to end, showing the interesting cases
- **A one-page architecture note**
- **A KPI summary** — baseline → expected → measured, plus a production measurement plan
- **A build trail** — how you built it with AI, as a **written document**. Keep it as you go.
- `results.json` — every decision your system made, validated against the schema

What's in the pack

`KICKOFF.md`	Dates, rules, environment setup, FAQ
`README.md`	The challenge brief
`current_state.docx`	The as-is process + your KPI baseline
`data/`	The data room — contracts, an email, a CSV
`policy_kb/`	The rules that decide each action
`schema/`	The record format your loads must match
`mock_erp/`	The live ERP API you build against
`selfcheck.py`	Run it now, and again before you submit
`SUBMISSION.md`	What to submit, how, and the deadline

Key dates

Mon 19 Oct	Kickoff — pack released, clock starts
Any time	Questions by email to the challenge mailbox
Fri 13 Nov	**Hard deadline, 23:59 US Eastern**
From 16 Nov	Judging: automated checks, panel review, finalist interviews
Early Dec	Winners announced across Digital Core

How you're scored

Pillar	Weight
AI-native SDLC — how you built it	25%
Agentic engineering — the architecture	30%
Structured outputs & data integrity	20%
Outcome & KPI reinvention	10%
Demo & communication	15%
Bonus — cost efficiency, eval harness, and handling the surprises in the data	**+10%**

Automated checks and scoring on every entry, then full panel review of a shortlist, then finalist interviews. No individual feedback — see SUBMISSION.md.

The rules

- **Individual entries.** Your build, end to end — finalist interviews confirm it.
- **Use any AI tooling you're licensed for.** Heavy AI use is the point, not a shortcut. Tooling is your own to source.
- **Don't touch the fixed integration.** The ERP app and the schemas are read-only and checksum-verified.
- **Synthetic data only.** Never bring real client or engagement data near this.
- **Questions by email only** — no office hours, and no design help. That's what's being scored.

What's in it for you

- **Performance points** for the top finishers
- **A fast track to FDE certification** — the rubric maps to the FDE competency model, and strong performers are flagged whether or not they place top 5
- **Visibility** — winners announced across Digital Core

Start here: unzip the pack and read `KICKOFF.md` (dates, rules, setup, FAQ), then `README.md` (the brief), then `current_state.docx` (the as-is process and your KPI baseline). `SUBMISSION.md` has the deadline and how to submit — read it in week one, not week four. Run `selfcheck.py` on day one to confirm your setup.